



Krutik K Mehta <krutikmehta@utexas.edu>

Fwd: Agentic Unlearning

1 message

Xiang Li · s <xl5@fordham.edu>

Thu, Jun 4, 2026 at 9:45 AM

To: "Mehta, Krutik" <kkm2563@my.utexas.edu>, "AlDarmaki, Amna" <aa224374@my.utexas.edu>

Hi Kurtis,

I forwarded the email I sent to Amna previously. Feel free to read these paper to get an idea. And feel free to reach out to me and Amna if you have any questions.

best,
Xiang

----- Forwarded message -----

From: **Xiang Li · s** <xl5@fordham.edu>

Date: Fri, May 8, 2026 at 03:22

Subject: Agentic Unlearning

To: <aa224374@my.utexas.edu>

Hi Amna,

Nice to meet you today. Below are some papers I think would be a good starting point.

For agentic unlearning, the two most relevant papers at the moment (to the best I know).

1. <https://arxiv.org/pdf/2602.17692>2. <https://arxiv.org/pdf/2604.00430>

These should help you understand the overall goal of agentic unlearning and how it differs from traditional LLM unlearning.

For Agentic AI topics, two closely related research topics are skill distillation and agentic memory. There's a little bit of overlap between them, but knowing them could be beneficial for you to know how agentic AI is built and what components there are for agent systems.

Skills

3. <https://arxiv.org/pdf/2603.25158>4. <https://arxiv.org/pdf/2602.08234>

Agentic memory

5. <https://arxiv.org/pdf/2502.12110>6. <https://arxiv.org/pdf/2602.13530>

If you are new to these areas, I'd also recommend checking some of the papers cited in their related work sections based on your interest.

I would suggest that you

1. Read these papers first to get the overall understanding of motivation, problem setting, and what people are currently trying to solve.

2. Then, start paying closer attention to the technical details: how the methods are designed, how memory/distillation is implemented, what datasets and evaluation metrics they use (very important to familiarize yourself with the datasets too)

3. If code is available, try running it and see if you can reproduce some of the results. It would be very helpful for understanding how the datasets are processed and how the pipeline is implemented because we might later modify or build on top of these frameworks once we decide on a concrete direction.

If you have time, it would also be great if you could prepare a few slides summarizing what you learned. That would help everyone get a clearer picture during the discussion, especially regarding the background, current methods, and datasets.

Also, feel free to reach out to me if you have any questions or want to discuss ideas!

best,
Xiang